

Algorithms for graph problems

Depth-First Search (DFS)

Here is a comprehensive guide to the Depth-First Search (DFS) algorithm, designed to be visual and easy to grasp without using complex mathematical notation.

What is Depth-First Search (DFS)?

Imagine you are in a massive maze. You want to find the exit. You stand at a fork in the road.

- **Strategy A:** You look down the left path, then the right path, just peeking a few steps in. (This is Breadth-First Search).
- **Strategy B (DFS):** You pick the left path and keep walking. If you hit another fork, you pick left again. You keep going until you hit a dead end. Only then do you retrace your steps (backtrack) to the last fork and try the right path.

Depth-First Search is Strategy B. It explores as far as possible along each branch before backtracking. It is an algorithm used to traverse or search tree and graph data structures.

Why Do We Need It?

DFS is fundamental for solving problems where you need to visit every location or find a specific path. Common uses include:

1. **Solving Puzzles:** Mazes, Sudoku, and crosswords (finding a solution by trying one path until it fails).
 2. **Finding Connected Components:** Checking if every computer in a network is connected to every other computer.
 3. **Pathfinding:** “Is there a path from Point A to Point B?” (Not necessarily the shortest one, but *a* path).
 4. **Cycle Detection:** Checking if a series of dependencies loops back on itself (like A depends on B, B depends on C, and C depends on A).
-

Visual Walkthrough: How It Works

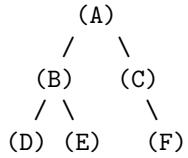
Let's trace a DFS on a simple graph.

The Rules:

1. Visit a node.
2. Mark it as “Visited” (so you don't go in circles).

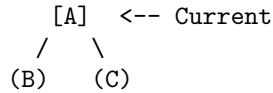
3. Go to an unvisited neighbor and repeat.
4. If stuck (no unvisited neighbors), **Backtrack** (return to the previous node).

The Graph:



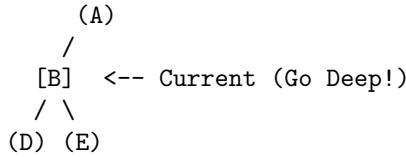
Step 1: Start at **A**. Mark A as visited.

- *Current Node:* A
- *Stack (History):* [A]




Step 2: A has neighbors B and C. Let's pick **B** (convention usually goes left first).

- *Current Node:* B
- *Stack (History):* [A, B]



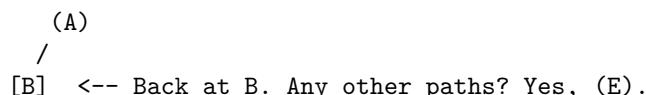
Step 3: B has neighbors D and E. We pick **D**.

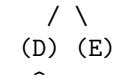
- *Current Node:* D
- *Stack (History):* [A, B, D]



Step 4: D has no neighbors (Dead End!). We **Backtrack** by popping D off the stack and returning to B.

- *Current Node:* B
- *Stack (History):* [A, B]

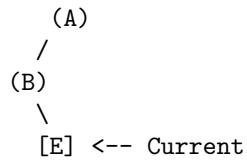




(Visited)

Step 5: From B, we visit the remaining neighbor **E**.

- *Current Node:* E
- *Stack (History):* [A, B, E]



Step 6: E is a dead end. Backtrack to B. B has no more unvisited neighbors. Backtrack to A.

- *Current Node:* A
 - *Stack (History):* [A]
- [A] <-- Back at A. Any other paths? Yes, (C).

(All visited)

Step 7: Visit **C**. Then Visit **F**.

- Final Path Traveled: A -> B -> D -> E -> C -> F.
-

Complexity Analysis

We measure algorithms by Time (how long it takes) and Space (how much memory it uses).

Time Complexity: $O(V + E)$

- **V** = Vertices (Nodes/Dots)
- **E** = Edges (Lines/Connections)

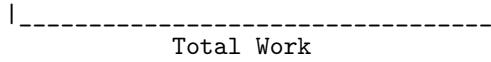
Why? Because in the worst case, you visit every Node once and travel across every Edge once.

Time Taken Visualization:

```

| 
| (Visiting every dot) + (Checking every line)
| 
| [ VVVVVVV ] + [ EEEEEEEEEE ]

```



Space Complexity: $O(V)$ Why? Because of the **Stack** (or Recursion). In the worst-case scenario (a long straight line), your recursion stack will hold all the vertices at once before it can backtrack.

Space Used (The Stack) in Worst Case:

Graph: A -- B -- C -- D -- E

Memory Stack:

```
| [E] | < Top
| [D] |
| [C] |
| [B] |
| [A] | < Bottom
+-----+
```

Width is small, Height = Number of Vertices (V)

Algorithm Code

Here is how you write it. The **Recursive** method is the most common and easiest to write for interviews.

Python (Recursive)

```
# Graph represented as a dictionary (Adjacency List)
graph = {
    'A': ['B', 'C'],
    'B': ['D', 'E'],
    'C': ['F'],
    'D': [],
    'E': [],
    'F': []
}

visited = set() # To keep track of visited nodes

def dfs_recursive(node):
    # 1. Base Case: If already visited, stop.
    if node in visited:
        return

    # 2. Mark the current node as visited
    visited.add(node)

    # Recur for all the neighbors of this vertex
    for neighbor in graph[node]:
        dfs_recursive(neighbor)
```

```

print(f"Visiting: {node}")
visited.add(node)

# 3. Explore neighbors
for neighbor in graph[node]:
    dfs_recursive(neighbor)

# Start the process
print("DFS Path:")
dfs_recursive('A')

```

JavaScript (Recursive)

```

const graph = {
    'A': ['B', 'C'],
    'B': ['D', 'E'],
    'C': ['F'],
    'D': [],
    'E': [],
    'F': []
};

const visited = new Set();

function dfsRecursive(node) {
    // 1. If already visited, stop
    if (visited.has(node)) {
        return;
    }

    // 2. Mark and Process
    console.log(`Visiting: ${node}`);
    visited.add(node);

    // 3. Explore neighbors
    const neighbors = graph[node];
    for (let neighbor of neighbors) {
        dfsRecursive(neighbor);
    }
}

console.log("DFS Path:");
dfsRecursive('A');

```

Thinking in LeetCode: “The Flood Fill Strategy”

When you see a LeetCode problem involving a **Grid (Matrix)** or **Islands**, think DFS immediately.

Example Scenario: “Number of Islands” (Given a grid of ‘1’s (land) and ‘0’s (water), count the islands).

Mental Model: Imagine you are looking at a map. You drop a bucket of paint on a piece of land. The paint flows to all connected land but stops at water.

1. **Iterate:** Loop through every cell in the grid.
2. **Trigger:** If you see a “1” (Land) that hasn’t been visited, increase your **IslandCount** by 1.
3. **DFS (The Paint):** Immediately trigger DFS on that cell.
 - The DFS turns that cell into a “0” (or marks it visited).
 - It then looks Up, Down, Left, Right.
 - It moves to any neighbor that is a “1” and calls DFS on them.
4. **Result:** When the DFS finishes, it has “erased” that entire island. The loop continues to find the next island.

```
Grid before DFS:      Grid after DFS starts at (0,0):  
1 1 0 0              x x 0 0  
1 1 0 0      ----> x x 0 0  
0 0 1 0              0 0 1 0  (The 'x's are marked visited)
```

Comparison: DFS vs. BFS

| Feature | DFS (Depth-First) | BFS (Breadth-First) |
|-----------------------|--|--|
| Concept | Maze Solver (boldly go forward) | Radar / Ripple (scan everywhere nearby) |
| Data Structure | Stack (LIFO - Last In, First Out) | Queue (FIFO - First In, First Out) |
| Best For | Puzzles, Connecting components | Finding the Shortest Path |
| Memory | Good ($O(h)$ where h is height) | Heavy ($O(w)$ where w is width/widest part) |

Final Tip: Use DFS when you need to search everything or just check for existence. Use BFS if you specifically need the *shortest* path in an unweighted graph.

Breadth-First Search (BFS)

Breadth-First Search (BFS) is a fundamental graph traversal algorithm. Think of it like a ripple in a pond: it starts at a center point and expands outward in concentric circles, visiting everything close by before moving further away.

1. The Problem It Solves

BFS is primarily used for finding the **shortest path** in an unweighted graph (a graph where the distance between all connected nodes is equal).

If you are looking for a key in a house, a random search might take you to the basement, then the attic, then the kitchen. BFS checks the room you are in, then the adjacent rooms, then the rooms adjacent to those. It guarantees that if a solution exists, you will find the one closest to the start.

Common scenarios:

- Finding the shortest route on a map (with equal road lengths).
 - Social networks: “Who are the friends of my friends?” (2nd degree connections).
 - Web crawlers: Analyzing links on a website level-by-level.
-

2. What is BFS?


BFS explores a graph **layer by layer**.

1. Start at a chosen node (Level 0).
2. Visit all immediate neighbors (Level 1).
3. Visit all neighbors of those neighbors (Level 2).
4. Repeat until all reachable nodes are visited.

It uses a **Queue** data structure (First-In, First-Out) to keep track of which node to visit next.

3. Visual Walkthrough

Let's look at this simple graph:



Step 1: Initialization We start at A. We put A in the Queue and mark it as “Visited”.

Queue: [A]
Visited: {A}

Step 2: Processing A We take A out of the queue. We look at its neighbors (B and C). We add them to the queue.

Processing: A
Queue: [B, C] <-- B and C are added to the back
Visited: {A, B, C}

Step 3: Processing B We take B out. We look at its neighbor (D). We add D to the queue.

Processing: B
Queue: [C, D] <-- C was already there, D joins the back
Visited: {A, B, C, D}

Step 4: Processing C We take C out. We look at its neighbor (E). We add E to the queue.

Processing: C
Queue: [D, E]
Visited: {A, B, C, D, E}

Step 5: Processing D & E We take D out. It has no unvisited neighbors. We take E out. It has no unvisited neighbors.

Queue: []
Status: Done.
Traversal Order: A → B → C → D → E

4. Complexity Analysis

- **V:** Vertices (Nodes)
- **E:** Edges (Connections)

Time Complexity: $O(V + E)$ We visit every node once $O(V)$ and for every node, we check its connections $O(E)$. Since we process the whole graph, the time taken grows linearly with the size of the graph.

Space Complexity: $O(V)$ In the worst case, the Queue might have to store the majority of the nodes at once (imagine a “star” graph where the center node is connected to 1,000 other nodes; after processing the center, all 1,000 neighbors go into the queue).

Worst Case Space (Star Graph)
[Center]

```
/ / | \
[1] [2] [3] [4] [5] ...
```

Queue becomes: [1, 2, 3, 4, 5...] (Huge width)

5. Code Implementation

Python

```
from collections import deque

def bfs(graph, start_node):
    visited = set()
    queue = deque([start_node])

    # Mark start as visited immediately
    visited.add(start_node)

    result = []

    while queue:
        # Dequeue from the front (FIFO)
        current_node = queue.popleft()
        result.append(current_node)

        # Check all neighbors
        for neighbor in graph[current_node]:
            if neighbor not in visited:
                visited.add(neighbor)
                queue.append(neighbor)

    return result

# Example Graph
adj_list = {
    'A': ['B', 'C'],
    'B': ['D'],
    'C': ['E'],
    'D': [],
    'E': []
}

print(bfs(adj_list, 'A'))
# Output: ['A', 'B', 'C', 'D', 'E']
```

JavaScript

```
function bfs(graph, startNode) {
    let visited = new Set();
    let queue = [startNode];
    visited.add(startNode);

    let result = [];

    while (queue.length > 0) {
        // Dequeue from front
        let currentNode = queue.shift();
        result.push(currentNode);

        let neighbors = graph[currentNode];
        for (let neighbor of neighbors) {
            if (!visited.has(neighbor)) {
                visited.add(neighbor);
                queue.push(neighbor);
            }
        }
    }
    return result;
}

const adjList = {
    'A': ['B', 'C'],
    'B': ['D'],
    'C': ['E'],
    'D': [],
    'E': []
};

console.log(bfs(adjList, 'A'));
// Output: ['A', 'B', 'C', 'D', 'E']
```

6. LeetCode Mental Model

When solving problems, look for these triggers to use BFS:

1. **“Shortest Path” / “Minimum Steps”:** If the problem asks for the minimum number of steps to convert X to Y (e.g., Word Ladder), use BFS.
2. **“Level Order”:** Anything requiring you to process a tree row-by-row.
3. **“Rotting Oranges” / “Virus Spread”:** Problems where a state

spreads to neighbors simultaneously each turn.

Strategy for “Grid” Problems (Islands, Mazes): Instead of an adjacency list, your “neighbors” are just the cells $(\text{row}+1, \text{col})$, $(\text{row}-1, \text{col})$, $(\text{row}, \text{col}+1)$, and $(\text{row}, \text{col}-1)$.

Grid Navigation Visual:

```
[ ] [ ] [ ]  
[ ] [X] [ ]  <- X is current  
[ ] [ ] [ ]
```

Neighbors are Up, Down, Left, Right.

Check bounds (don't go off grid) and check visited set.

7. Comparison with Other Algorithms

| Feature | BFS (Breadth-First) | DFS (Depth-First) |
|-----------------------|--|--|
| Data Structure | Queue (FIFO) | Stack (LIFO) or Recursion |
| Visual Style | Expands in a circle (Wave) | Snakes deep into one path |
| Shortest Path? | Yes (Guaranteed for unweighted) | No (Might take a long detour) |
| Memory | High (Stores width of “wave”) | Low (Stores length of path) |
| Best For | Finding closest target | Mazes, Puzzles requiring all solutions |

Analogy:

- **BFS:** Scanning a crowd of people by looking at everyone in the front row, then everyone in the second row.
- **DFS:** Walking up to one person, asking who they know, walking to that person, asking who *they* know, until you hit a dead end.

Multi-Source Breadth First Search (Multi-Source BFS)

Multi-Source Breadth-First Search (BFS)

1. The Problem: Why Do We Need It?

Imagine you drop a single stone into a calm pond. Ripples expand outward from that one point. This is **Standard BFS**. It tells you the shortest distance

from *one* starting point to anywhere else.

Now, imagine you drop **three** stones into the pond at the exact same time at different locations. Ripples expand from *all three* points simultaneously. If you want to know how soon a ripple reaches a specific leaf floating on the water, you care about the ripple that gets there *first*.

The Problem: You have a graph (or grid) and a set of “source” nodes. You need to find the shortest distance from *any* source node to every other node in the graph.

If you ran a Standard BFS from Source A, then another from Source B, and another from Source C, and then compared the results, it would be incredibly slow and inefficient. **Multi-Source BFS** solves this in a single pass.

2. What Is It?

Multi-Source BFS is a slight modification of the standard BFS algorithm.

- **Standard BFS:** Initialize the Queue with **one** starting node with distance 0.
- **Multi-Source BFS:** Initialize the Queue with **all** starting nodes with distance 0 simultaneously.

Functionally, you can imagine a “fake” invisible node (a Super Source) connected to all your actual starting nodes with an edge weight of 0. You run standard BFS from this Super Source.

3. Visual Walkthrough

Let’s visualize a **virus spreading** on a grid.

- S = Source (Infected)
- . = Empty (Uninfected)
- X = Wall (Blockage)

The Grid:

| | | | |
|-------|-------|-------|-------|
| (0,0) | (0,1) | (0,2) | (0,3) |
| S | . | S | . |
| . | . | . | . |
| . | X | . | . |

Step 0: Initialization We identify all S nodes. We add them to our Queue. We set their distance to 0. All other nodes are “unvisited” (infinity).

- **Queue:** [(0,0), (0,2)]
- **Visited/Distance Map:**

| | | | |
|-----|-----|-----|-----|
| 0 | inf | 0 | inf |
| inf | inf | inf | inf |
| inf | inf | inf | inf |

Step 1: Process Distance 0 Nodes We pop $(0,0)$ and $(0,2)$ one by one. We look at their neighbors (Up, Down, Left, Right). If a neighbor is valid and not visited, we infect it (set distance to 1) and add it to the queue.

- Pop $(0,0)$: Neighbors are $(0,1)$ and $(1,0)$.
- Pop $(0,2)$: Neighbors are $(0,1)$ (already seen via 0,0? No, let's say we process 0,0 first, so 0,1 gets dist 1), $(0,3)$, $(1,2)$.
- **Queue:** $[(0,1), (1,0), (0,3), (1,2)]$ (These are the “1”s)
- **Grid State (Current Distances):**

| | | | |
|-----|-----|-----|-----|
| 0 | 1 | 0 | 1 |
| 1 | inf | 1 | inf |
| inf | X | inf | inf |

Step 2: Process Distance 1 Nodes The “ripples” expand again. We process the nodes currently in the queue. Their unvisited neighbors become distance 2.

- Pop $(0,1)$: Neighbors $(1,1)$ becomes 2.
- Pop $(1,0)$: Neighbor $(2,0)$ becomes 2.
- Pop $(0,3)$: Neighbor $(1,3)$ becomes 2.
- Pop $(1,2)$: Neighbors $(1,1)$ (already visited), $(1,3)$ (already visited), $(2,2)$ becomes 2.
- **Queue:** $[(1,1), (2,0), (1,3), (2,2)]$
- **Grid State:**

| | | | |
|---|---|---|-----|
| 0 | 1 | 0 | 1 |
| 1 | 2 | 1 | 2 |
| 2 | X | 2 | inf |

Step 3: Process Distance 2 Nodes Final expansion.

- Pop $(1,1)$: Neighbors $(1,0)$ (visited), $(1,2)$ (visited), $(0,1)$ (visited), $(2,1)$ (It's a wall ‘X’, ignore).
- ...
- Pop $(2,2)$: Neighbor $(2,3)$ becomes 3.
- **Final Grid:**

| | | | |
|---|---|---|---|
| 0 | 1 | 0 | 1 |
| 1 | 2 | 1 | 2 |
| 2 | X | 2 | 3 |

This tells us exactly how many steps it takes for the virus to reach any cell, regardless of which source it came from.

4. Implementation

Python Code Using `collections.deque` for an efficient O(1) pop from the left.

```
from collections import deque

def multi_source_bfs(grid):
    rows = len(grid)
    cols = len(grid[0])

    # Queue for BFS
    queue = deque()

    # Visited set or distance matrix initialized to infinity
    # Using a distance matrix allows us to track steps
    dist = [[-1 for _ in range(cols)] for _ in range(rows)]

    # 1. INITIALIZATION: Add ALL sources to the queue
    for r in range(rows):
        for c in range(cols):
            if grid[r][c] == 'S': # 'S' is a source
                queue.append((r, c))
                dist[r][c] = 0

    # Directions: Up, Down, Left, Right
    directions = [(-1, 0), (1, 0), (0, -1), (0, 1)]

    # 2. STANDARD BFS LOOP
    while queue:
        curr_r, curr_c = queue.popleft()

        current_distance = dist[curr_r][curr_c]

        for dr, dc in directions:
            new_r, new_c = curr_r + dr, curr_c + dc

            # Check bounds
            if 0 <= new_r < rows and 0 <= new_c < cols:
                # Check if not visited (dist is -1) and not a wall
                if dist[new_r][new_c] == -1 and grid[new_r][new_c] != 'X':
                    dist[new_r][new_c] = current_distance + 1
                    queue.append((new_r, new_c))

    return dist
```

JavaScript Code

```
function multiSourceBFS(grid) {
    const rows = grid.length;
    const cols = grid[0].length;
    const queue = [];

    // Initialize distance matrix with -1 (representing infinity/unvisited)
    const dist = Array.from({ length: rows }, () => Array(cols).fill(-1));

    // 1. INITIALIZATION
    for (let r = 0; r < rows; r++) {
        for (let c = 0; c < cols; c++) {
            if (grid[r][c] === 'S') {
                queue.push([r, c]);
                dist[r][c] = 0;
            }
        }
    }

    const directions = [[-1, 0], [1, 0], [0, -1], [0, 1]];
    let head = 0; // Optimization for queue pointer rather than shifting array

    // 2. BFS LOOP
    while (head < queue.length) {
        const [currR, currC] = queue[head++];
        const currentDist = dist[currR][currC];

        for (const [dr, dc] of directions) {
            const newR = currR + dr;
            const newC = currC + dc;

            if (newR >= 0 && newR < rows && newC >= 0 && newC < cols) {
                if (dist[newR][newC] === -1 && grid[newR][newC] !== 'X') {
                    dist[newR][newC] = currentDist + 1;
                    queue.push([newR, newC]);
                }
            }
        }
    }
    return dist;
}
```

5. Complexity Analysis

Let V be the number of vertices (or cells $R * C$) and E be the number of edges.

Time Complexity: $O(V + E)$

- Why? Even though we have multiple sources, every node is added to the queue and processed **at most once**.
- In a grid, E is roughly $4 * V$, so it simplifies to $O(V)$ or $O(R * C)$.

Space Complexity: $O(V)$

- Why? In the worst case (e.g., a completely filled grid), the queue might hold a significant portion of the nodes. We also store the `dist` array which is size V .
- In a grid: $O(R * C)$.

Visualization of Complexity

| Node Processing (Time) | Memory Usage (Space) |
|-------------------------------------|-------------------------------------|
| [X] Processed Once | [Q] Queue Storage |
| [X] Processed Once | [D] Distance Matrix |
| [X] Processed Once | |
| ... | |
| Total: Linear Prop. to Grid Size | Total: Linear Prop. to Grid Size |

6. LeetCode Strategy: How to Think

When you see a problem, ask yourself these questions to decide if you need Multi-Source BFS:

1. **“Nearest X to Y”:** Does the problem ask for the distance to the *nearest* specific item (like nearest gate, nearest zero, nearest safe zone)?
2. **Simultaneous Start:** Does the process happen everywhere at once? (e.g., rotten oranges rotting neighbors simultaneously).
3. **Parallel Flooding:** Can I visualize the problem as water flooding from multiple leaks?

Classic LeetCode Examples:

- **01 Matrix (LeetCode 542):** Find distance of nearest 0 for each cell.
- *Strategy:* Push all ‘0’ cells into queue first.
- **Rotting Oranges (LeetCode 994):** Minimum time until all oranges rot.
- *Strategy:* Push all initially rotten oranges into queue.

- **Map of Highest Peak (LeetCode 1765):** Assign heights such that water cells are 0 and height diff is at most 1.
 - *Strategy:* Push all water cells into queue.
-

7. Comparison with Other Algorithms

| Feature | Single-Source BFS | Multi-Source BFS | Dijkstra | Floyd-Warshall |
|----------------------|----------------------|-------------------------------|--------------------------|------------------|
| Input Sources | One (e.g., Start) | Many (e.g., All Monsters) | One | All Pairs |
| Weights | Unweighted (1) | Unweighted (1) | Weighted (non-negative) | Weighted |
| Best For | Shortest path A to B | Shortest path {Set A} to rest | Shortest path with costs | All-to-All paths |
| Speed | Fast O(V+E) | Fast O(V+E) | Slower O(E log V) | Very Slow O(V^3) |

Key Takeaway: Multi-Source BFS is just BFS, but you “cheat” by starting the race with multiple runners already on the track.

Disjoint Set Union: Union-Find Algorithm

Here is a detailed, easy-to-understand guide to the Disjoint Set Union (DSU), also known as the Union-Find algorithm.

1. Why Do We Need It? (The Problem)

Imagine you are managing a social network. You have a massive list of people, and occasionally, two people become “friends.” Friends of friends are also considered connected.

The Challenge: You frequently get two types of queries:

1. **Connect:** “Alice and Bob just became friends.”
2. **Query:** “Are Alice and Charlie connected (directly or indirectly)?”

If you use standard Graph algorithms like BFS or DFS (Breadth-First Search) to answer “Are they connected?”, it takes **O(N)** time (linear time) for *every* query. If you have millions of users and millions of queries, BFS is too slow.

The Solution: We need a way to connect items and check connectivity almost instantly. This is what Union-Find does. It solves the **Dynamic Connectivity Problem**.

2. What Is It?

Union-Find is a data structure that keeps track of a set of elements partitioned into a number of disjoint (non-overlapping) subsets.

Think of it as a collection of “gangs” or “clubs.”

- Initially, everyone is in their own club of one.
- **Union:** Merging two clubs together.
- **Find:** Asking “Which club does this person belong to?” (represented by the club’s leader).

If Person A and Person B have the same leader, they are in the same club.

3. How It Works

We represent these sets as **Trees**.

- Each element is a node.
- Each subset is a tree.
- The “Leader” (or Representative) is the **root** of the tree.

Key Optimizations To make this algorithm lightning fast, we use two tricks:

1. **Path Compression (The “Shortcut”):** When we look for the leader of a node, we make that node point *directly* to the leader so we don’t have to traverse the whole chain next time. **Before Find(D):** A (Leader) -> B -> C -> D **After Find(D):** A (Leader) -> B A (Leader) -> C A (Leader) -> D (*Now C and D point directly to A*)
 2. **Union by Rank/Size (The “Balance”):** When merging two trees, we always attach the **shorter** tree under the **taller** tree. This prevents the tree from becoming a long, thin line (linked list), keeping it flat and fast.
-

4. Visual Walkthrough

Let’s say we have 5 elements: 0, 1, 2, 3, 4.

Step 1: Initialization Everyone is their own parent (leader).

[0] [1] [2] [3] [4]
 ^ ^ ^ ^ ^

| | | | |
Leader Leader Leader Leader Leader

- Parent Array: [0, 1, 2, 3, 4]

Step 2: Union(0, 1) Connect 0 and 1. We make 0 the leader of 1.

0 (Leader)
/
1

[2] [3] [4]

- Parent Array: [0, 0, 2, 3, 4]

Step 3: Union(2, 3) Connect 2 and 3. We make 2 the leader of 3.

0 2 (Leader)
/
1 3

[4]

- Parent Array: [0, 0, 2, 2, 4]

Step 4: Union(1, 4) We find the leader of 1 (which is 0). We find the leader of 4 (which is 4). We attach 4 to 0.

0 (Leader) 2 (Leader)
/ \ /
1 4 3

- Parent Array: [0, 0, 2, 2, 0]

Step 5: Union(3, 4) This connects the two big components.

1. Find leader of 3 -> **2**
2. Find leader of 4 -> **0**
3. Attach 2 (shorter tree) under 0 (taller tree).

0 (Ultimate Leader)
/ | \
1 4 2
|
3

- Now, **Find(3)** will return 0. **Find(1)** returns 0. They are connected.
-

5. Complexity Analysis

Using both Path Compression and Union by Rank, the time complexity is nearly constant.

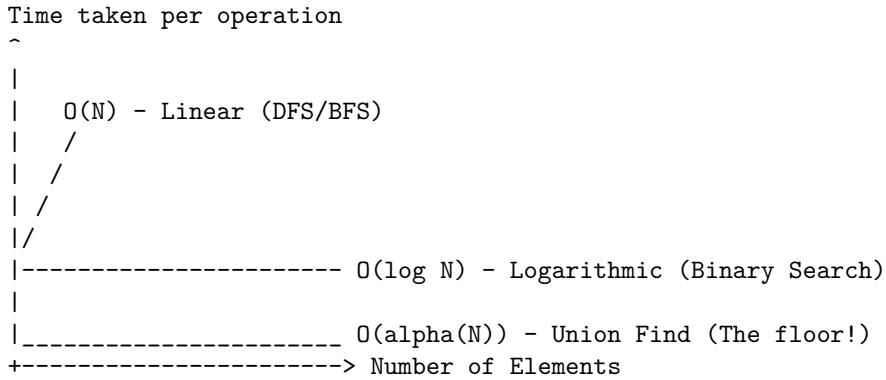
Time Complexity: $O(\alpha(N))$

- α is the **Inverse Ackermann function**.
- For all practical values of N (even up to the number of atoms in the universe), $\alpha(N)$ is **less than 5**.
- It is essentially $O(1)$ (constant time) for practical purposes.

Space Complexity: $O(N)$

- We need an array to store the parent of each node.

Visualization of Speed:



6. Implementation

Here is the efficient implementation with Path Compression and Union by Rank.

Python

```
class UnionFind:  
    def __init__(self, size):  
        # Initially, each node is its own parent  
        self.parent = list(range(size))  
        # Rank helps to keep the tree flat (height approximation)  
        self.rank = [1] * size  
  
    def find(self, p):  
        # Path Compression:  
        # If p is not its own parent, recursively find the root  
        # and update p's parent to directly point to the root.  
        if self.parent[p] != p:  
            self.parent[p] = self.find(self.parent[p])  
        return self.parent[p]
```

```

    if self.parent[p] != p:
        self.parent[p] = self.find(self.parent[p])
    return self.parent[p]

def union(self, p, q):
    rootP = self.find(p)
    rootQ = self.find(q)

    # If they are already in the same set, do nothing
    if rootP == rootQ:
        return False

    # Union by Rank: Attach smaller tree to larger tree
    if self.rank[rootP] > self.rank[rootQ]:
        self.parent[rootQ] = rootP
    elif self.rank[rootP] < self.rank[rootQ]:
        self.parent[rootP] = rootQ
    else:
        self.parent[rootQ] = rootP
        self.rank[rootP] += 1

    return True

# Usage
dsu = UnionFind(5)
dsu.union(0, 1)
dsu.union(2, 3)
dsu.union(1, 4)
print(dsu.find(3) == dsu.find(4)) # False (before connecting components)
dsu.union(3, 4)
print(dsu.find(3) == dsu.find(4)) # True (after connecting)

```

JavaScript

```

class UnionFind {
    constructor(size) {
        // Create array [0, 1, 2, ..., size-1]
        this.parent = Array.from({ length: size }, (_, i) => i);
        this.rank = new Array(size).fill(1);
    }

    find(p) {
        // Path Compression
        if (this.parent[p] !== p) {
            this.parent[p] = this.find(this.parent[p]);
        }
        return this.parent[p];
    }
}

```

```

        return this.parent[p];
    }

    union(p, q) {
        const rootP = this.find(p);
        const rootQ = this.find(q);

        if (rootP === rootQ) return false;

        // Union by Rank
        if (this.rank[rootP] > this.rank[rootQ]) {
            this.parent[rootQ] = rootP;
        } else if (this.rank[rootP] < this.rank[rootQ]) {
            this.parent[rootP] = rootQ;
        } else {
            this.parent[rootQ] = rootP;
            this.rank[rootP] += 1;
        }
        return true;
    }
}

```

7. How to Apply to LeetCode Problems

Problem Idea: “Redundant Connection” (LeetCode 684)

- **The Prompt:** You are given a set of edges. A graph is a tree if it has no cycles. Find an edge that, if removed, makes the graph a tree (i.e., find the edge that created a cycle).

Strategy: Iterate through every edge $[u, v]$ in the input.

1. Check $\text{find}(u)$ and $\text{find}(v)$.
2. If $\text{find}(u) == \text{find}(v)$, it means u and v are **already connected** by a previous path. Adding this edge $[u, v]$ would create a loop (a cycle). **This is your answer!**
3. If they are not connected, $\text{union}(u, v)$ and continue.

Other Identifiers: Use Union-Find when you see keywords like:

- “Number of connected components”
 - “Group the items together”
 - “Equivalence classes”
 - “Graph validity” (checking for cycles)
-

8. Comparison with Other Algorithms

| Feature | Union-Find | DFS / BFS |
|------------------------|---|---|
| Primary Use | Dynamic Connectivity (adding edges live) | Static Connectivity (graph is fixed) |
| Handling Cycles | Excellent (Detects cycle instantly) | Good (Detects via visited set) |
| Path Finding | Cannot find the <i>path</i> (only if connected) | Can find the actual path between nodes |
| Complexity | Almost $O(1)$ per operation | $O(V + E)$ per traversal |
| Implementation | Array-based (Very compact) | Stack/Queue + Adjacency List |

Single Source Shortest Path: Dijkstra's Algorithm

Here is a comprehensive guide to **Dijkstra's Algorithm**, explained using plain text, ASCII visualizations, and code, strictly adhering to your preference for non-mathematical notation.

1. The Problem: Why Do We Need This?

Imagine you are driving and using a GPS. You are at **Point A** (Home) and want to get to **Point B** (Work).

- There are many roads (edges) connecting many intersections (nodes).
- Each road has a “cost” or “weight.” This could be **distance**, **time** (traffic), or **tolls**.

You don't just want *any* path; you want the **cheapest** path (shortest distance or fastest time).

The Challenge: A simple breadth-first search (BFS) assumes every road takes the same amount of time (weight = 1). In the real world, highways are faster than dirt roads. We need an algorithm that respects these different weights.

2. What is Dijkstra's Algorithm?

Dijkstra's Algorithm finds the shortest path from a single starting point (Source) to **all other points** in a graph with **non-negative weights**.

Key Analogy: Imagine dumping a bucket of water on the starting node. The water spreads along the pipes (edges). It will reach closer nodes first before flowing to further ones. Dijkstra's algorithm calculates exactly when the water reaches every other node.

Core Principles:

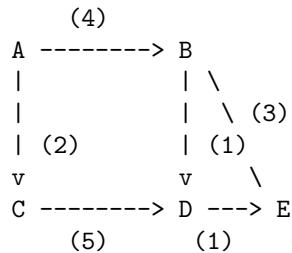
1. **Greedy Approach:** Always process the closest “unvisited” node next.
 2. **Relaxation:** If we find a shortcut to a node through the current node, we update its distance.
- *Example:* If I thought it took 10 minutes to get to Node X, but I found a path through Node Y that takes only 8 minutes, I “relax” the edge and update Node X’s distance to 8.
-

3. Visual Walkthrough (ASCII)

Let’s find the shortest path from **Start Node (A)** to all other nodes.

The Graph:

- Nodes: A, B, C, D, E
- Weights are numbers on the lines.



Connections (Weights):

- A -> B (4)
- A -> C (2)
- C -> B (1) <- Note: This creates a shortcut to B!
- C -> D (5)
- B -> D (1)
- B -> E (3)
- D -> E (1)

Initialization

We track two things:

1. **Shortest Distance:** Known distance from A to a node (starts at ‘infinity’ for everyone except A).

2. **Priority Queue:** A list of nodes to visit, sorted by shortest distance.

Start State:

- Current Node: **A** (Distance: 0)
- Unvisited: {B, C, D, E}

DISTANCES:

```
[ A: 0 ]
[ B: inf ]
[ C: inf ]
[ D: inf ]
[ E: inf ]
```

Step 1: Visit Node A

Look at neighbors of A: **B** and **C**.

- Distance to B = A's dist (0) + weight (4) = **4**. (4 is less than inf, so update).
- Distance to C = A's dist (0) + weight (2) = **2**. (2 is less than inf, so update).

Mark A as visited. Who is next? We look at the unvisited nodes with the smallest distance: B(4) vs C(2). **C is smaller.**

STATUS UPDATE:

```
Visited: {A}
Queue: [ C: 2, B: 4 ]  <-- C is at the front because 2 < 4
```

DISTANCES:

```
[ A: 0 ]
[ B: 4 ]
[ C: 2 ]  <-- Next to visit
[ D: inf ]
[ E: inf ]
```

Step 2: Visit Node C (Distance: 2)

Neighbors of C: **B** and **D**.

- **Check B:**
- Current known distance to B is **4** (via A->B).
- New path via C: (Distance to C) + (C->B) = 2 + 1 = **3**.
- **3 < 4**, so we found a shortcut! **Update B to 3.**

- **Check D:**

- New path via C: (Distance to C) + (C->D) = 2 + 5 = **7**.
- Update D to 7.

Mark C as visited. Who is next? Smallest in Queue: B(3) vs D(7). **B is smaller.**

STATUS UPDATE:

Visited: {A, C}

Queue: [B: 3, D: 7]

DISTANCES:

```
[ A: 0 ]
[ B: 3 ] <-- Updated from 4! (Shortcut found)
[ C: 2 ]
[ D: 7 ]
[ E: inf ]
```

Step 3: Visit Node B (Distance: 3)

Neighbors of B: **D** and **E**.

- **Check D:**

- Current known distance to D is **7** (via C).
- New path via B: (Distance to B) + (B->D) = 3 + 1 = **4**.
- **4 < 7**, shortcut found! **Update D to 4**.

- **Check E:**

- New path via B: (Distance to B) + (B->E) = 3 + 3 = **6**.
- Update E to 6.

Mark B as visited. Who is next? Smallest in Queue: D(4) vs E(6). **D is smaller.**

STATUS UPDATE:

Visited: {A, C, B}

Queue: [D: 4, E: 6]

DISTANCES:

```
[ A: 0 ]
[ B: 3 ]
[ C: 2 ]
[ D: 4 ] <-- Updated from 7!
[ E: 6 ]
```

Step 4: Visit Node D (Distance: 4)

Neighbors of D: E.

- **Check E:**
- Current known distance to E is **6**.
- New path via D: (Distance to D) + (D->E) = 4 + 1 = **5**.
- **5 < 6**, shortcut found! **Update E to 5**.

Mark D as visited. Only E is left.

STATUS UPDATE:

Visited: {A, C, B, D}

Queue: [E: 5]

DISTANCES:

[A: 0]
[B: 3]
[C: 2]
[D: 4]
[E: 5] <-- Updated from 6!

Final Result

Shortest paths from A:

- A: 0
 - B: 3 (Path: A -> C -> B)
 - C: 2 (Path: A -> C)
 - D: 4 (Path: A -> C -> B -> D)
 - E: 5 (Path: A -> C -> B -> D -> E)
-

4. Complexity Analysis (ASCII)

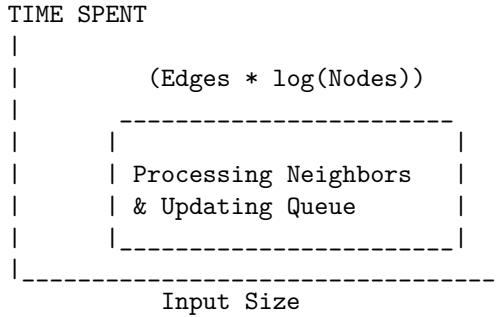
Here is how the algorithm scales.

- **V** = Number of Vertices (Nodes)
- **E** = Number of Edges (Connections)

Time Complexity

Using a Min-Priority Queue (Heap): **O(E * log V)**

Think of it as: For every road (Edge), we might have to sort the list of cities ($\log V$).



Space Complexity

$O(V + E)$ We need to store the graph ($V + E$) and the distances/priority queue (V).

MEMORY USED

| [Graph Storage] | [Distances Array] |
|-----------------|-------------------|
| (Nodes + Edges) | (Nodes) |
| $O(V + E)$ | $O(V)$ |

5. Code Implementation

Python

We use `heapq` because it is efficient.

```

import heapq

def dijkstra(graph, start_node):
    # 1. Initialize distances to infinity, start_node to 0
    # Dictionary to store shortest distance to each node
    distances = {node: float('inf') for node in graph}
    distances[start_node] = 0

    # 2. Priority Queue: stores tuples of (current_distance, node)
    # We start with the source node
    priority_queue = [(0, start_node)]

    while priority_queue:
        # Pop the node with the smallest distance
        current_dist, current_node = heapq.heappop(priority_queue)

        # Optimization: If we found a shorter way to this node already, skip

```

```

    if current_dist > distances[current_node]:
        continue

    # Explore neighbors
    for neighbor, weight in graph[current_node].items():
        distance = current_dist + weight

        # If a shorter path is found
        if distance < distances[neighbor]:
            distances[neighbor] = distance
            # Push to queue with new priority
            heapq.heappush(priority_queue, (distance, neighbor))

    return distances

# Example Graph (from the ASCII diagrams)
graph = {
    'A': {'B': 4, 'C': 2},
    'B': {'D': 1, 'E': 3},
    'C': {'B': 1, 'D': 5},
    'D': {'E': 1},
    'E': {}
}

print(dijkstra(graph, 'A'))
# Output: {'A': 0, 'B': 3, 'C': 2, 'D': 4, 'E': 5}

```

JavaScript

JavaScript does not have a built-in Priority Queue, so for simple problems, we can use an array and sort it (less efficient) or implement a MinHeap class (efficient). Below is a conceptual version using a simple array sort for readability.

```

function dijkstra(graph, startNode) {
    // 1. Initialize distances
    let distances = {};
    for (let node in graph) {
        distances[node] = Infinity;
    }
    distances[startNode] = 0;

    // 2. Priority Queue (Simple Array for demo)
    // In production, use a proper MinHeap data structure
    let pq = [[0, startNode]];

    while (pq.length > 0) {

```

```

// Sort to simulate Priority Queue (Smallest distance first)
pq.sort((a, b) => a[0] - b[0]);

// Get node with smallest distance
let [currentDist, currentNode] = pq.shift();

if (currentDist > distances[currentNode]) continue;

// Check neighbors
let neighbors = graph[currentNode];
for (let neighbor in neighbors) {
    let weight = neighbors[neighbor];
    let newDist = currentDist + weight;

    if (newDist < distances[neighbor]) {
        distances[neighbor] = newDist;
        pq.push([newDist, neighbor]);
    }
}
return distances;
}

const graph = {
    'A': {'B': 4, 'C': 2},
    'B': {'D': 1, 'E': 3},
    'C': {'B': 1, 'D': 5},
    'D': {'E': 1},
    'E': {}
};

console.log(dijkstra(graph, 'A'));

```

6. LeetCode Strategy: How to use this

When you see a problem on LeetCode, ask yourself these questions to see if Dijkstra is the right tool:

1. **Is it a graph?** (Are there states/nodes and connections?)
2. **Are there weights?** (Does moving from A to B cost something different than C to D?)
3. **Are weights non-negative?** (Dijkstra fails if weights are negative).
4. **Are you looking for a minimum cost/time/path?**

Mental Template for Solving:

“I need to track the ‘best so far’ cost for every node. I will use a Min-Heap to always process the cheapest option next. If I find a cheaper way to a node, I update it and add it to the heap.”


Common LeetCode variations:

- **Network Delay Time:** Find how long it takes for a signal to reach all nodes. (Classic Dijkstra).
 - **Cheapest Flights Within K Stops:** Dijkstra, but with a constraint on the number of edges used (often requires a slight modification or Bellman-Ford if straightforward).
 - **Path with Minimum Effort:** The “weight” might be the absolute difference in height between two cells in a grid.
-

7. Comparison with Other Algorithms

| Algorithm | Best Used For | Key Difference | Weights Allowed |
|---------------------------------------|----------------------------------|--|----------------------------|
| BFS (Breadth-First Search) | Unweighted Graphs | Explores layer by layer. Like Dijkstra but assumes all edge weights are 1. | N/A (assumes 1) |
| Dijkstra | Weighted Graphs | Uses a Priority Queue to prioritize lower costs. | Must be Positive |
| Bellman-Ford | Graphs with Negative Weights | Can detect negative cycles. Much slower than Dijkstra. | Positive & Negative |
| A* (A-Star) | Pathfinding to a specific Target | Uses a “Heuristic” (guess) to guide the search toward the goal faster. | Positive |

Summary Algorithm Choice Flowchart (ASCII)



```

    v
Are there Negative Weights?
|
YES ----> Use Bellman-Ford
|
NO
|
    v
Dijkstra's Algorithm (Min-Heap)

```

Single Source Shortest Path: Bellman Ford Algorithm

Here is a detailed guide to the Bellman-Ford Algorithm, designed to be clear and visual without using complex mathematical notation.

1. The Problem: Why do we need Bellman-Ford?

In graph theory, the **Single Source Shortest Path (SSSP)** problem asks: “*If I start at Node A, what is the shortest path to every other node in the graph?*”

Most people use **Dijkstra’s Algorithm** for this. However, Dijkstra’s algorithm fails if the graph has **negative edge weights**.

- **Example:** Imagine a game where traveling usually costs energy (positive weight), but sliding down a magic chute *gives* you energy back (negative weight). Dijkstra cannot handle the magic chute correctly.

The Solution: The **Bellman-Ford Algorithm** can calculate shortest paths even if edges have negative weights. It can also detect **Negative Weight Cycles** (infinite loops where you keep gaining energy forever).

2. What is it and How does it work?

Bellman-Ford is based on the **Relaxation Principle**. It is slower than Dijkstra but more versatile.

The Core Idea: “Relaxation”

To “relax” an edge means to check if we can improve the shortest path to a node by going through a specific neighbor.

- Current best distance to Node B is infinity.
- Current best distance to Node A is 5.

- There is an edge from A \rightarrow B with weight 2.
- **Relaxation:** Is $(\text{Distance to A}) + (\text{Weight A} \rightarrow \text{B}) < (\text{Current Distance to B})$?
- Is $5 + 2 < \text{infinity}$?
- Yes! Update distance to Node B to 7.

The Algorithm Logic

1. **Initialize:** Set the distance to the Source Node to 0 and all other nodes to Infinity.
 2. **Iterate:** Repeat the following process $V - 1$ times (where V is the number of vertices/nodes):
 - Look at **every single edge** in the graph.
 - Try to relax that edge.
 3. **Check for Cycles:** Run the relaxation one more time. If any distance still changes, it means there is a **Negative Weight Cycle**.
-

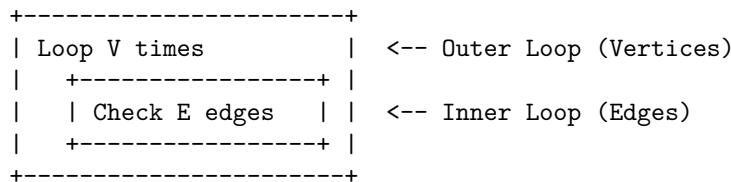
3. Complexity Analysis

Here is the breakdown of the efficiency using ASCII visualizations.

Time Complexity

Since we iterate $(V-1)$ times and check all E edges each time:

Time Complexity: $O(V * E)$



- **Best Case:** $O(E)$ (If the edges happen to be ordered perfectly, though standard implementation is always $O(V*E)$).
- **Worst Case:** $O(V * E)$.

Space Complexity

We only need to store the distance to each node.

Space Complexity: $O(V)$

[Dist A | Dist B | Dist C | ... | Dist V]

4. Visual Walkthrough (ASCII)


Let's find the shortest path from **Node A**.

The Graph:

- Nodes: A, B, C
- Edges:
- A -> B (weight: 4)
- A -> C (weight: 2)
- C -> B (weight: -10) (**Negative Edge!**)

Step 0: Initialization

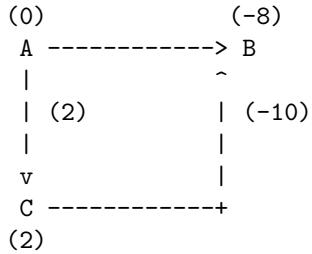
Set Source (A) to 0, others to Infinity (Inf).



Iteration 1 (Relax all edges)

1. Check A -> B (weight 4):
 - Is $\text{Dist}(A) + 4 < \text{Dist}(B)$?
 - $0 + 4 < \text{Inf}$? YES.
 - Update B to 4.
2. Check A -> C (weight 2):
 - Is $\text{Dist}(A) + 2 < \text{Dist}(C)$?
 - $0 + 2 < \text{Inf}$? YES.
 - Update C to 2.
3. Check C -> B (weight -10):
 - Is $\text{Dist}(C) + (-10) < \text{Dist}(B)$?
 - $2 + (-10) < 4$?
 - $-8 < 4$? YES.
 - Update B to -8.

End of Iteration 1:



Note: We normally repeat this $V-1$ times (2 times here). In the second iteration, no values would change because the paths are already optimal.

Final Result:

- Path A \rightarrow B is cost -8 (Route: A \rightarrow C \rightarrow B).
 - Path A \rightarrow C is cost 2.
-

5. Implementation

Python Code

```

def bellman_ford(graph, num_nodes, start_node):
    # Step 1: Initialize distances
    # Using float('inf') to represent infinity
    distances = {node: float('inf') for node in range(num_nodes)}
    distances[start_node] = 0

    # Step 2: Relax edges V-1 times
    # 'graph' is a list of tuples: (u, v, weight)
    for _ in range(num_nodes - 1):
        for u, v, weight in graph:
            if distances[u] != float('inf') and distances[u] + weight < distances[v]:
                distances[v] = distances[u] + weight

    # Step 3: Check for Negative Cycles
    for u, v, weight in graph:
        if distances[u] != float('inf') and distances[u] + weight < distances[v]:
            print("Graph contains negative weight cycle")
            return None

    return distances

# Example Usage
# Nodes: 0, 1, 2
# Edges: 0->1 (4), 0->2 (2), 2->1 (-10)
edges = [

```

```

        (0, 1, 4),
        (0, 2, 2),
        (2, 1, -10)
    ]
print(bellman_ford(edges, 3, 0))
# Output: {0: 0, 1: -8, 2: 2}

```

JavaScript Code

```

function bellmanFord(edges, numNodes, startNode) {
    // Step 1: Initialize distances
    let distances = new Array(numNodes).fill(Infinity);
    distances[startNode] = 0;

    // Step 2: Relax edges V-1 times
    for (let i = 0; i < numNodes - 1; i++) {
        for (let [u, v, weight] of edges) {
            if (distances[u] !== Infinity && distances[u] + weight < distances[v]) {
                distances[v] = distances[u] + weight;
            }
        }
    }

    // Step 3: Check for Negative Cycles
    for (let [u, v, weight] of edges) {
        if (distances[u] !== Infinity && distances[u] + weight < distances[v]) {
            console.log("Graph contains negative weight cycle");
            return null;
        }
    }

    return distances;
}

// Example Usage
const edges = [
    [0, 1, 4],
    [0, 2, 2],
    [2, 1, -10]
];
console.log(bellmanFord(edges, 3, 0));
// Output: [0, -8, 2]

```

6. How to Use this for LeetCode

When you see a graph problem on LeetCode, use this mental checklist to decide if Bellman-Ford is the right tool:

1. Does the problem mention costs/weights that can be negative?
 - Yes: Immediate Bellman-Ford signal.
2. Does the problem ask to find a “cycle” that reduces cost?
 - Yes: This is a Negative Cycle Detection problem.
3. Is the graph small?
 - Bellman-Ford is $O(V^2E)$. If $V > 5,000$, this might Time Out (TLE). Dijkstra is faster ($O(E \log V)$) but stricter on weights.

Specific LeetCode Problem Strategy:

- **Problem:** “Cheapest Flights Within K Stops”
 - **Strategy:** This is a modified Bellman-Ford. Instead of running the loop $V-1$ times, you run it exactly $K+1$ times. This restricts the path to only finding routes that use K stops or fewer.
-

7. Comparison Table

| Feature | Bellman-Ford | Dijkstra | BFS (Breadth-First Search) |
|------------------------|--|-----------------------------|--------------------------------------|
| Edge Weights | Handles Positive & Negative | Must be Non-Negative | Unweighted (all edges = 1) |
| Time Complexity | Slow: $O(V * E)$ | Fast: $O(E \log V)$ | Very Fast: $O(V + E)$ |
| detects Cycles? | Yes (Negative Cycles) | No | No (for weights) |
| Use Case | Financial graphs (arbitrage), small graphs with weird rules. | GPS, Maps, Network routing. | Peer-to-peer networks, web crawlers. |

Practice Problem with negative weights: “The Time Travel Paradox”

The Scenario: Imagine you are navigating a network of star systems. Most routes take time (positive cost). However, there is a “wormhole” that actually sends you back in time (negative cost).

The Question: Determine if there is a path in this star map where you can keep looping forever to travel infinitely far back in the past. In Computer Science terms: “**Does this graph contain a Negative Weight Cycle?**”

1. The Graph Setup (ASCII)

- **Nodes (Star Systems):** 0, 1, 2
- **Edges (Routes):**
- 0 -> 1 (Cost: 1)
- 1 -> 2 (Cost: -5) <- *The Wormhole*
- 2 -> 0 (Cost: 3)

Cycle Math: If you go $0 \rightarrow 1 \rightarrow 2 \rightarrow 0$: $1 + (-5) + 3 = -1$ Every time you complete this loop, the total cost decreases by 1. This is a **Negative Cycle**.

2. Step-by-Step Algorithm Walkthrough

We have **3 Nodes** ($V=3$). We must run the relaxation loop **V-1 times** (2 times). Finally, we run it **1 extra time** to detect the cycle.

Initialization: Distances: [0, Inf, Inf] (Start at Node 0)

Iteration 1 (Standard Bellman-Ford Pass) We check all edges:

1. **Edge 0->1 (weight 1):**
 - $\text{Dist}[0] + 1 < \text{Dist}[1]$? $-> 0 + 1 < \text{Inf}$? **Yes.**
 - Update $\text{Dist}[1]$ to 1.
 - *State:* [0, 1, Inf]
2. **Edge 1->2 (weight -5):**
 - $\text{Dist}[1] + (-5) < \text{Dist}[2]$? $-> 1 - 5 < \text{Inf}$? **Yes.**
 - Update $\text{Dist}[2]$ to -4.
 - *State:* [0, 1, -4]
3. **Edge 2->0 (weight 3):**
 - $\text{Dist}[2] + 3 < \text{Dist}[0]$? $-> -4 + 3 < 0$? **Yes.**
 - Update $\text{Dist}[0]$ to -1.
 - *State:* [-1, 1, -4]

Iteration 2 (Standard Bellman-Ford Pass) We check all edges again using the new values [-1, 1, -4]:

1. **Edge 0->1 (weight 1):**
 - $\text{Dist}[0] + 1 < \text{Dist}[1]$? $-> -1 + 1 < 1$? **Yes.**

- Update $\text{Dist}[1]$ to **0**.
 - *State: $[-1, 0, -4]$*
2. Edge **1->2 (weight -5)**:
- $\text{Dist}[1] - 5 < \text{Dist}[2]?$ $\rightarrow 0 - 5 < -4?$ **Yes.**
 - Update $\text{Dist}[2]$ to **-5**.
 - *State: $[-1, 0, -5]$*
3. Edge **2->0 (weight 3)**:
- $\text{Dist}[2] + 3 < \text{Dist}[0]?$ $\rightarrow -5 + 3 < -1?$ **Yes.**
 - Update $\text{Dist}[0]$ to **-2**.
 - *State: $[-2, 0, -5]$*

The Detection Step (Iteration 3) According to the algorithm, after $V-1$ iterations, we should be done. If we can still relax an edge, a cycle exists.

Let's test **Edge 0->1** again:

- Current $\text{Dist}[0]$ is **-2**.
- Current $\text{Dist}[1]$ is **0**.
- Check: $\text{Dist}[0] + 1 < \text{Dist}[1]?$
- $-2 + 1 < 0?$ $\rightarrow -1 < 0?$ **YES!**

Conclusion: Because we could still reduce the cost in the extra step, we have proven a **Negative Weight Cycle** exists. Infinite time travel is possible!

3. Solution Code

Here is how you would write a function specifically to **detect** this boolean condition.

Python Solution

```
def has_negative_cycle(num_nodes, edges, start_node):
    # Step 1: Initialize
    dist = {i: float('inf') for i in range(num_nodes)}
    dist[start_node] = 0

    # Step 2: Relax V-1 times
    for _ in range(num_nodes - 1):
        for u, v, weight in edges:
            if dist[u] != float('inf') and dist[u] + weight < dist[v]:
                dist[v] = dist[u] + weight

    # Step 3: Cycle Detection Check
    # We iterate through all edges ONE more time
```

```

        for u, v, weight in edges:
            if dist[u] != float('inf') and dist[u] + weight < dist[v]:
                print(f"Cycle detected! Node {u} to {v} can still be reduced.")
                return True # Cycle exists!

        return False # Safe graph

# Test with our Time Travel graph
# Nodes: 0, 1, 2
# Edges: (0->1, w=1), (1->2, w=-5), (2->0, w=3)
my_edges = [
    (0, 1, 1),
    (1, 2, -5),
    (2, 0, 3)
]

print(f"Negative Cycle Detected? {has_negative_cycle(3, my_edges, 0)}")

```

JavaScript Solution

```

function hasNegativeCycle(numNodes, edges, startNode) {
    let dist = new Array(numNodes).fill(Infinity);
    dist[startNode] = 0;

    // Step 2: Relax V-1 times
    for (let i = 0; i < numNodes - 1; i++) {
        for (let [u, v, weight] of edges) {
            if (dist[u] !== Infinity && dist[u] + weight < dist[v]) {
                dist[v] = dist[u] + weight;
            }
        }
    }

    // Step 3: Cycle Detection Check
    for (let [u, v, weight] of edges) {
        if (dist[u] !== Infinity && dist[u] + weight < dist[v]) {
            console.log(`Cycle detected between ${u} and ${v}`);
            return true; // Cycle exists!
        }
    }

    return false;
}

const myEdges = [
    [0, 1, 1],

```

```

[1, 2, -5],
[2, 0, 3]
];

console.log("Negative Cycle Detected?", hasNegativeCycle(3, myEdges, 0));

```

Minimum Spanning Tree (MST): Prim's Algorithm

What Problem Does It Solve?

Imagine you are a city planner who needs to connect several neighborhoods with fiber optic cables.

- **The Goal:** Every neighborhood must be connected to the network (directly or indirectly).
- **The Constraint:** You want to minimize the total cost of the cable used.
- **The Catch:** You don't need redundant loops. If Neighborhood A is connected to B, and B to C, you don't need a direct cable from A to C if the cost is higher.

This is the **Minimum Spanning Tree (MST)** problem. You have a “graph” of points (vertices) and possible connections (edges) with costs (weights). You need to select a subset of edges that connects everyone with the lowest total weight and no cycles.

What is Prim's Algorithm?

Prim's Algorithm is a “greedy” strategy to find this MST. It builds the tree one node at a time, always picking the cheapest connection available from the nodes you have *already visited* to a node you *haven't visited yet*.

Core Logic:

1. Start at an arbitrary node.
 2. Look at all edges connecting your current “tree” to the outside world.
 3. Pick the shortest (cheapest) edge that leads to a new node.
 4. Add that node and edge to your tree.
 5. Repeat until all nodes are included.
-

Visual Walkthrough

Let's trace Prim's algorithm on a simple graph.

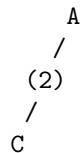
The Graph: Nodes: A, B, C, D, E

Weights:

- A-B: 4
- A-C: 2
- B-C: 3
- B-D: 2
- B-E: 3
- C-D: 4
- C-E: 5
- D-E: 1

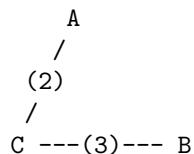
Step 1: Start at Node A

- **Visited:** {A}
- **Available Edges from Visited:**
 - A -> B (cost 4)
 - A -> C (cost 2)
- **Decision:** A -> C is cheaper (2 vs 4). Choose A-C.



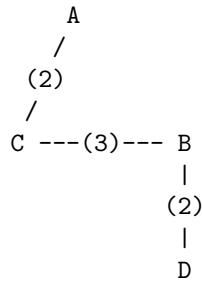
Step 2: Tree is now {A, C}

- **Visited:** {A, C}
- **Available Edges from {A, C} to Unvisited:**
 - A -> B (cost 4)
 - C -> B (cost 3)
 - C -> D (cost 4)
 - C -> E (cost 5)
- **Decision:** C -> B is the cheapest (cost 3). Choose C-B.
- **Note:** We ignore A->B (cost 4) because we found a cheaper way to get to B via C.




Step 3: Tree is now {A, C, B}

- **Visited:** {A, C, B}
- **Available Edges from {A, C, B} to Unvisited:**
- C -> D (cost 4)
- C -> E (cost 5)
- B -> D (cost 2)
- B -> E (cost 3)
- **Decision:** B -> D is the cheapest (cost 2). Choose B-D.



Step 4: Tree is now {A, C, B, D}

- **Visited:** {A, C, B, D}
- **Available Edges from {A, C, B, D} to Unvisited:**
- C -> E (cost 5)
- B -> E (cost 3)
- D -> E (cost 1)
- **Decision:** D -> E is the cheapest (cost 1). Choose D-E.



Result: All nodes {A, B, C, D, E} are visited. Total Cost: $2 + 3 + 2 + 1 = 8$

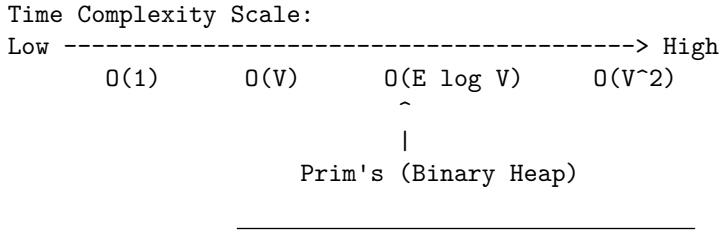
Complexity Analysis

Here is the breakdown of efficiency using a Min-Priority Queue (the standard implementation).

V = Number of Vertices (Nodes) E = Number of Edges

| Metric | Complexity | Explanation |
|--------------|-----------------|--|
| Time | $O(E * \log V)$ | For every edge, we might push/pop from the priority queue. The heap operations take logarithmic time relative to the number of vertices. |
| Space | $O(V + E)$ | We need to store the graph (adjacency list) and the arrays to track weights and visited status. |

Ascii Visualization of Complexity:



Implementation

Here is how you would code this efficiently using a Min-Heap (Priority Queue).

Python Code

```
import heapq

def prims_algorithm(n, edges):
    # n: number of nodes
    # edges: list of [u, v, weight]

    # Build Adjacency List
    adj = {i: [] for i in range(n)}
    for u, v, w in edges:
```

```

adj[u].append((v, w))
adj[v].append((u, w)) # MST is for undirected graphs

# Min-Heap to pick smallest weight edge
# Format: (weight, node_to)
min_heap = [(0, 0)] # Start at node 0 with cost 0

visited = set()
total_cost = 0
mst_edges = 0

while min_heap:
    weight, u = heapq.heappop(min_heap)

    # If we already visited this node, skip it
    if u in visited:
        continue

    visited.add(u)
    total_cost += weight
    mst_edges += 1

    # If we have visited all nodes, we can stop early
    if mst_edges == n:
        break

    # Check neighbors
    for v, w in adj[u]:
        if v not in visited:
            heapq.heappush(min_heap, (w, v))

return total_cost if len(visited) == n else -1

# Example Usage
# Nodes 0, 1, 2, 3, 4 represent A, B, C, D, E
edges_data = [
    [0, 1, 4], [0, 2, 2], [1, 2, 3],
    [1, 3, 2], [1, 4, 3], [2, 3, 4],
    [2, 4, 5], [3, 4, 1]
]
print(prims_algorithm(5, edges_data)) # Output: 8

```

JavaScript Code

```

class MinPriorityQueue {
    // Simple implementation of a Min Priority Queue for visualization

```

```

constructor() { this.heap = []; }

enqueue(element, priority) {
    this.heap.push({ element, priority });
    this.heap.sort((a, b) => a.priority - b.priority); // O(N log N) for simplicity here
}

dequeue() { return this.heap.shift(); }
isEmpty() { return this.heap.length === 0; }
}

function primsAlgorithm(n, edges) {
    const adj = new Map();
    for (let i = 0; i < n; i++) adj.set(i, []);

    for (const [u, v, w] of edges) {
        adj.get(u).push([v, w]);
        adj.get(v).push([u, w]);
    }

    const pq = new MinPriorityQueue();
    pq.enqueue(0, 0); // Start at node 0, cost 0

    const visited = new Set();
    let totalCost = 0;

    while (!pq.isEmpty()) {
        const { element: u, priority: weight } = pq.dequeue();

        if (visited.has(u)) continue;

        visited.add(u);
        totalCost += weight;

        if (visited.size === n) break;

        for (const [v, w] of adj.get(u)) {
            if (!visited.has(v)) {
                pq.enqueue(v, w);
            }
        }
    }

    return visited.size === n ? totalCost : -1;
}

```

```
// Example Usage
const edgesData = [
  [0, 1, 4], [0, 2, 2], [1, 2, 3],
  [1, 3, 2], [1, 4, 3], [2, 3, 4],
  [2, 4, 5], [3, 4, 1]
];
console.log(primsAlgorithm(5, edgesData)); // Output: 8
```

Algorithm Comparisons

It is helpful to know when to use Prim's versus others.

| Feature | Prim's Algorithm | Kruskal's Algorithm | Dijkstra's Algorithm |
|-----------------------|---|---|--|
| Goal | Minimum Spanning Tree (MST) | Minimum Spanning Tree (MST) | Shortest Path from Source |
| Strategy | Grow a single tree from a source node. | Sort all edges by weight and merge disjoint sets. | Grow path from source, minimizing total distance from start. |
| Best For | Dense Graphs (lots of edges per node). | Sparse Graphs (fewer edges). | Finding shortest route (GPS). |
| Data Structure | Priority Queue + Visited Set | Union-Find (Disjoint Set) + Sorting | Priority Queue + Distance Array |
| Analogy | A mold spreading to nearest food. | Building bridges between islands until all connected. | Driving directions. |

Applying to LeetCode: “Min Cost to Connect All Points”

A classic problem is **LeetCode 1584: Min Cost to Connect All Points**.

The Setup: You are given an array of coordinates `points = [[x1,y1], [x2,y2]...]`. You need to connect all points. The cost is the Manhattan distance: $|x_1 - x_2| + |y_1 - y_2|$.

How to Think About It:

1. **Identify the structure:** “Connect all points” + “Min Cost” = Minimum Spanning Tree.
2. **Dense Graph:** Since *every* point can connect to *every other* point, this is a “Complete Graph” (Dense).
3. **Selection:** Prim’s is usually preferred over Kruskal’s for dense graphs because Kruskal’s requires sorting $O(E \log E)$ edges. In a complete graph, E is roughly V^2 , making sorting very expensive. Prim’s avoids processing every single edge explicitly if implemented optimally.

Strategy:

- Treat point 0 as your starting node.
- Calculate the distance from point 0 to all others and push to a Min-Heap.
- Pop the closest point, add it to visited, and calculate distances from *that* point to remaining unvisited points.

Would you like me to generate a solution specifically for that LeetCode problem?

Minimum Spanning Tree (MST): Kruskal’s Algorithm

Here is a detailed guide to Minimum Spanning Trees (MST) and Kruskal’s Algorithm.

The Problem: Why Do We Need This?

Imagine you are a city planner who needs to connect several neighborhoods with fiber-optic cables.

- **Goal:** You must ensure *every* neighborhood is connected to the network.
- **Constraint:** You have a limited budget, so you want to use the least amount of cable possible.
- **Rule:** You don’t need redundant loops (cycles). If Neighborhood A connects to B, and B connects to C, A is already connected to C. Adding a direct cable from A to C is unnecessary expense.

This is the **Minimum Spanning Tree** problem. You need a subset of connections (edges) that connects all points (vertices) with the minimum total cost (weight).

What is Kruskal’s Algorithm?

Kruskal’s Algorithm is a **greedy** strategy to solve the MST problem. It builds the spanning tree by adding edges one by one into a growing spanning tree.

The Strategy:

1. **Be Cheap:** Always pick the smallest weight edge available in the entire graph.
2. **Be Safe:** Only add that edge if it doesn't create a loop (cycle). If it creates a loop, discard it.
3. **Repeat:** Keep doing this until all vertices are connected.

To handle the “cycle detection” efficiently, Kruskal’s uses a data structure called **Disjoint Set Union (DSU)** or **Union-Find**. This keeps track of which vertices are already in the same group.

Visual Walkthrough

Let’s visualize this with a graph.

The Graph:

- **Vertices (Nodes):** 0, 1, 2, 3
- **Edges (Connections):**
- 0 to 1 (Weight: 10)
- 0 to 2 (Weight: 6)
- 0 to 3 (Weight: 5)
- 1 to 3 (Weight: 15)
- 2 to 3 (Weight: 4)

Step 1: Sort All Edges First, we list all edges and sort them from smallest weight to largest.

| Source | Destination | Weight | |
|--------|-------------|--------|--------------|
| 2 | 3 | 4 | <-- Smallest |
| 0 | 3 | 5 | |
| 0 | 2 | 6 | |
| 0 | 1 | 10 | |
| 1 | 3 | 15 | <-- Largest |

Step 2: Processing Edges

Iteration 1: Pick edge (2, 3) with weight 4.

- Are 2 and 3 connected? No.
- **Action:** Add edge. Join sets {2} and {3}.

Visual:

(0) (1)


(2)===== (3) <-- Edge added

Iteration 2: Pick edge (0, 3) with weight 5.

- Are 0 and 3 connected? No.

- **Action:** Add edge. Join set {0} with set {2, 3}.


Visual:



Iteration 3: Pick edge (0, 2) with weight 6.

- Are 0 and 2 connected?
- Trace: 0 is connected to 3, and 3 is connected to 2.
- Yes, they are in the same group.
- **Action:** Discard (SKIP). Adding this would create a triangle cycle (0-2-3).

Visual (No Change):




Attempted to add connection here (0-2), but rejected.

Iteration 4: Pick edge (0, 1) with weight 10.

- Are 0 and 1 connected? No.
- **Action:** Add edge. Join set {1} with set {0, 2, 3}.

Visual:



Iteration 5: Pick edge (1, 3) with weight 15.

- Are 1 and 3 connected? Yes (via 0).
- **Action:** Discard (SKIP).

Final Result:

- Edges: (2,3), (0,3), (0,1)
- Total Minimum Cost: $4 + 5 + 10 = \mathbf{19}$

Complexity Analysis

Here is the breakdown of the efficiency.

- V = Number of Vertices
- E = Number of Edges

| Type | Complexity | Explanation |
|-------|---------------|---|
| Time | $O(E \log E)$ | Sorting the edges takes $O(E \log E)$. The Union-Find operations take nearly constant time, usually written as $O(E * \alpha(V))$, which is very fast. The sorting is the bottleneck. |
| Space | $O(V + E)$ | We need space to store the list of edges $O(E)$ and the parent array for Union-Find $O(V)$. |

Implementation

We will use a `UnionFind` helper class to manage the groups efficiently.

Python

```
class UnionFind:  
    def __init__(self, n):  
        # Initially, each node is its own parent  
        self.parent = list(range(n))  
        # Rank helps keep the tree flat during union  
        self.rank = [0] * n  
  
    def find(self, i):  
        # Path compression: Point directly to the root  
        if self.parent[i] != i:  
            self.parent[i] = self.find(self.parent[i])  
        return self.parent[i]  
  
    def union(self, i, j):  
        root_i = self.find(i)  
        root_j = self.find(j)
```

```

    if root_i != root_j:
        # Union by rank: attach smaller tree to larger tree
        if self.rank[root_i] > self.rank[root_j]:
            self.parent[root_j] = root_i
        elif self.rank[root_i] < self.rank[root_j]:
            self.parent[root_i] = root_j
        else:
            self.parent[root_j] = root_i
            self.rank[root_i] += 1
        return True # Union was successful
    return False # Cycle detected

def kruskals(num_vertices, edges):
    """
    edges: List of [source, dest, weight]
    """
    mst_weight = 0
    mst_edges = []

    # Step 1: Sort edges by weight (ascending)
    # x[2] corresponds to the weight
    edges.sort(key=lambda x: x[2])

    uf = UnionFind(num_vertices)

    for u, v, w in edges:
        # Step 2: Check if adding edge creates a cycle
        if uf.union(u, v):
            mst_weight += w
            mst_edges.append([u, v, w])

    return mst_weight, mst_edges

# Example Usage
# Nodes: 0, 1, 2, 3
edges_list = [
    [0, 1, 10],
    [0, 2, 6],
    [0, 3, 5],
    [1, 3, 15],
    [2, 3, 4]
]

weight, tree = kruskals(4, edges_list)
print("MST Weight:", weight)
print("Edges in MST:", tree)

```

JavaScript

```
class UnionFind {
    constructor(n) {
        this.parent = Array.from({ length: n }, (_, i) => i);
        this.rank = new Array(n).fill(0);
    }

    find(i) {
        if (this.parent[i] !== i) {
            this.parent[i] = this.find(this.parent[i]);
        }
        return this.parent[i];
    }

    union(i, j) {
        let rootI = this.find(i);
        let rootJ = this.find(j);

        if (rootI !== rootJ) {
            if (this.rank[rootI] > this.rank[rootJ]) {
                this.parent[rootJ] = rootI;
            } else if (this.rank[rootI] < this.rank[rootJ]) {
                this.parent[rootI] = rootJ;
            } else {
                this.parent[rootJ] = rootI;
                this.rank[rootI] += 1;
            }
            return true; // Edge added
        }
        return false; // Cycle detected
    }
}

function kruskals(numVertices, edges) {
    let mstWeight = 0;
    let mstEdges = [];

    // Step 1: Sort edges by weight
    edges.sort((a, b) => a[2] - b[2]);

    const uf = new UnionFind(numVertices);

    // Step 2: Iterate and build MST
    for (let [u, v, w] of edges) {
        if (uf.union(u, v)) {
            mstWeight += w;
            mstEdges.push([u, v, w]);
        }
    }
}
```

```

        mstWeight += w;
        mstEdges.push([u, v, w]);
    }
}

return { mstWeight, mstEdges };
}

// Example Usage
const edges = [
    [0, 1, 10], [0, 2, 6], [0, 3, 5],
    [1, 3, 15], [2, 3, 4]
];

const result = kruskals(4, edges);
console.log("MST Weight:", result.mstWeight);
console.log("MST Edges:", result.mstEdges);

```

Strategy for LeetCode Problems

When you see a problem on LeetCode, use Kruskal's if you spot these clues:

1. **“Connect all points”:** The problem asks to link cities, computers, or nodes.
2. **“Minimum Cost”:** You need to minimize the sum of weights (cost, distance, time).
3. **Sparse Graphs:** Kruskal's is often easier to implement and slightly faster on graphs with fewer edges compared to vertices (Sparse Graphs).

Mental Model for Solving: Think of the problem as a “Merging Islands” scenario.

- Initially, every node is its own island.
- Bridges (edges) are available at different prices.
- Buy the cheapest bridges first.
- If a bridge connects two nodes that are already on the same island (via other bridges), don't buy it.

Example Problem: “Min Cost to Connect All Points”

- **Input:** Coordinates of points.
 - **Step 1:** Calculate the Manhattan distance between every pair of points (these are your edge weights).
 - **Step 2:** Sort all these distances.
 - **Step 3:** Run Kruskal's loop.
-

Comparison: Kruskal's vs. Prim's

Both algorithms solve the exact same problem, but they approach it differently.

| Feature | Kruskal's Algorithm | Prim's Algorithm |
|-----------------------|---|--|
| Approach | Edge-centric: Focuses on the edges. Sorts them and picks the best one anywhere in the graph. | Vertex-centric: Focuses on growing a single tree from a starting node. Grows like a mold spreading out. |
| Data Structure | Uses Union-Find (Disjoint Set). | Uses a Priority Queue (Min-Heap). |
| Best For | Sparse Graphs: (Lots of nodes, fewer edges). It is simpler to implement with an edge list. | Dense Graphs: (Lots of edges connecting everything). Since it doesn't need to sort <i>all</i> edges upfront, it can be faster here. |
| Connectivity | Can generate a “forest” (multiple disconnected trees) if the graph is disconnected. | Only generates a single tree from the start node. |

Topological Sort (Kahn's Algorithm)

Here is a detailed, easy-to-understand guide to Topological Sort using Kahn's Algorithm.

1. The Problem: The “Dependency” Nightmare

Imagine you are trying to bake a cake, or perhaps enroll in university courses. You cannot just do things in any random order.

- **Baking:** You must mix the batter *before* you put it in the oven.
- **College:** You must take “Intro to CS” *before* you take “Advanced Algorithms.”

This is the **Dependency Problem**. We have a set of tasks, and some tasks must be completed before others can start.

Topological Sort is the algorithm that takes a jumbled list of tasks and dependencies and straightens them out into a linear order where every prerequisite comes *before* the task it enables.

Key Rule: This only works on **Directed Acyclic Graphs (DAGs)**. * **Directed:** A leads to B (one way). * **Acylic:** No loops! If A waits for B, and B waits for A, you are stuck forever.

2. What is Kahn's Algorithm?

Kahn's Algorithm is a specific way to perform a Topological Sort. It relies on a very simple, intuitive concept called **Indegree**.

- **Indegree:** The number of arrows pointing *at* a node.
- **Meaning:** If a task has an Indegree of 0, it has **no prerequisites**. It is free to be done immediately!

The Strategy

1. Find all tasks with **0 dependencies** (Indegree = 0).
 2. “Do” those tasks (add them to our sorted list).
 3. Once a task is done, remove it from the graph.
 4. Check if this removal freed up any new tasks (reduced their Indegree to 0).
 5. Repeat until the graph is empty.
-

3. Visual Walkthrough


Let's trace this with a simple graph representing getting dressed.

The Dependencies:

- Shirt (A) -> Tie (C)
- Shirt (A) -> Jacket (D)
- Pants (B) -> Belt (E)
- Tie (C) -> Jacket (D)

Initial State

Nodes: A(Shirt), B(Pants), C(Tie), D(Jacket), E(Belt)



Step 1: Calculate Indegrees Count how many arrows point *into* each letter.


- A: 0 (No arrows pointing in)
- B: 0
- C: 1 (From A)
- D: 2 (From A and C)
- E: 1 (From B)

Queue: [A, B] (These have 0 indegree, ready to process) **Sorted List:** []

Iteration 1

Pop **A** from the Queue. Add to Sorted List. “Remove” A’s arrows. This reduces the indegree of its neighbors (C and D).

Processed: A




Updates:

- C was 1, now **0**. (Added to Queue!)
 - D was 2, now **1**.
 - Queue is now: [B, C]
 - Sorted List: [A]
-

Iteration 2

Pop **B** from the Queue. Add to Sorted List. “Remove” B’s arrows. This reduces the indegree of neighbor E.

Processed: B



Updates:

- E was 1, now **0**. (Added to Queue!)
 - Queue is now: [C, E]
 - Sorted List: [A, B]
-

Iteration 3

Pop **C** from the Queue. Add to Sorted List. “Remove” C’s arrows. Reduces indegree of D.

Processed: C



Updates:

- D was 1, now **0**. (Added to Queue!)
 - Queue is now: [E, D]
 - Sorted List: [A, B, C]
-

Iteration 4 & 5

Pop **E** (no neighbors). Pop **D** (no neighbors).

Final Sorted List: [A, B, C, E, D] (*Note: “Shirt, Pants, Tie, Belt, Jacket” is a valid order to get dressed!*)

4. Code Implementation

Here is the logic in Python and JavaScript.

Python

```
from collections import deque

def topological_sort(num_nodes, edges):
    # 1. Initialize Graph and Indegrees
    graph = {i: [] for i in range(num_nodes)}
    indegree = {i: 0 for i in range(num_nodes)}

    # 2. Build the graph
    for src, dest in edges:
        graph[src].append(dest)
        indegree[dest] += 1

    # 3. Add 0-indegree nodes to queue
    queue = deque()
    for node in range(num_nodes):
        if indegree[node] == 0:
            queue.append(node)

    sorted_order = []

    # 4. Process Queue
    while queue:
        current = queue.popleft()
        sorted_order.append(current)

        # Reduce indegree of neighbors
        for neighbor in graph[current]:
            indegree[neighbor] -= 1
            if indegree[neighbor] == 0:
                queue.append(neighbor)

    # Check for cycles (if sorted_order length < num_nodes)
    if len(sorted_order) != num_nodes:
        return [] # Cycle detected!

    return sorted_order

# Example Usage
# 0: Shirt, 1: Pants, 2: Tie, 3: Jacket, 4: Belt
print(topological_sort(5, [[0, 2], [0, 3], [1, 4], [2, 3]]))
```

JavaScript

```
function topologicalSort(numNodes, edges) {
    // 1. Initialize Graph and Indegrees
    const graph = Array.from({ length: numNodes }, () => []);
    const indegree = new Array(numNodes).fill(0);

    // 2. Build the graph
    for (const [src, dest] of edges) {
        graph[src].push(dest);
        indegree[dest]++;
    }

    // 3. Add 0-indegree nodes to queue
    const queue = [];
    for (let i = 0; i < numNodes; i++) {
        if (indegree[i] === 0) queue.push(i);
    }

    const sortedOrder = [];

    // 4. Process Queue
    while (queue.length > 0) {
        const current = queue.shift();
        sortedOrder.push(current);

        for (const neighbor of graph[current]) {
            indegree[neighbor]--;
            if (indegree[neighbor] === 0) {
                queue.push(neighbor);
            }
        }
    }

    // Cycle check
    return sortedOrder.length === numNodes ? sortedOrder : [];
}
```


5. Complexity Analysis

We define:

- **V** = Number of Vertices (Nodes/Tasks)
- **E** = Number of Edges (Dependencies/Arrows)


Time Complexity: $O(V + E)$

Why? We visit every node exactly once and we travel across every edge exactly once.



Space Complexity: $O(V + E)$

We need to store the graph (Adjacency List) and the indegree array.



6. LeetCode Application Strategy

Use this thought process for problems like **LeetCode 207 (Course Schedule)** or **210 (Course Schedule II)**.

The “Is it Kahn’s?” Checklist:

1. **Is there an order?** Does the problem ask for a “valid sequence,” “prerequisites,” or “scheduling”?
2. **Is it directed?** Does A imply B, but B doesn’t imply A?
3. **Is cycle detection needed?** Does the problem imply failure if there’s a loop (e.g., “return false if you can’t finish all courses”)?

Mental Framework: Think of the graph as a waterfall.

- Water flows from top (Indegree 0) to bottom.

- Kahn's algorithm simply peels off the dry layer at the top, letting the water flow down to the next layer.
 - If you peel everything off and there are still nodes left but no “Indegree 0” nodes, you have a **pool (Cycle)** where water is stuck circling.
-

7. Comparison: Kahn's vs. DFS

Both solve Topological Sort, but they feel different.

| Feature | Kahn's Algorithm (BFS based) | DFS Approach |
|------------------------|--|--|
| Method | Iterative. Peels nodes one by one. | Recursive. Dives deep, adds to list as it backtracks. |
| Direction | Builds the list naturally: [First \rightarrow Last] | Builds list in reverse: [Last \rightarrow First], then reverses it. |
| Cycle Detection | Very Easy. If <code>result_length < v</code> , there is a cycle. | Tricky. Requires a recursion stack tracker (<code>visiting</code> vs <code>visited</code> sets). |
| Intuition | “Do what I can do <i>now</i> .” | “Find the end, then work backwards.” |

Trie Data Structure

Here is a detailed guide to the Trie (pronounced “try” or “tree”) data structure.

1. The Problem: Why do we need a Trie?

Imagine you are building the “Autocomplete” feature for a search engine. You have a database of 1 million words.

If a user types “ap”, you want to instantly suggest: “apple”, “apply”, “appetite”.

The Naive Approach (List of Strings): If you store all words in a simple list or array, you have to loop through every single word to check if it starts with “ap”.

- Word 1: “banana” (No)
- Word 2: “canary” (No)
- ...
- Word 500,000: “apple” (Yes!)

This is too slow. If you have millions of words, the search lags.

The Solution (Trie): A Trie organizes words by their *characters*. Instead of checking every word, you just follow the path: a -> p -> p ...

It solves the problem of **efficient prefix matching**.

2. What is a Trie?

A Trie (often called a **Prefix Tree**) is a tree-based data structure used to store a dynamic set of strings.

Key Characteristics:

- **Nodes:** Each node represents a single character.
 - **Root:** The top node is empty.
 - **Edges:** Links connect characters to form words.
 - **End Marker:** We usually mark the end of a valid word (often with a boolean flag or a special node).
-


3. How It Works: A Visual Walkthrough

Let's build a Trie containing these three words: "CAT", "CAR", "DO".

Step 1: The Empty Root We start with just a root node.


(ROOT)

Step 2: Insert "CAT" We check if 'C' exists. No? Add it. Then 'A'. Then 'T'. Mark 'T' as the end of a word (*).



Step 3: Insert "CAR" Start at Root.


- Do we have 'C'? Yes. Move there.
- Do we have 'A'? Yes. Move there.
- Do we have 'R'? No. Add it. Mark 'R' as end (*).



Notice how “CAT” and “CAR” share the path “C-A”. This saves space!

Step 4: Insert “DO” Start at Root.

- Do we have ‘D’? No. Add it (new branch).
- Add ‘O’. Mark as end (*).



4. Time and Space Complexity

Let’s look at the costs using L as the length of the word you are searching/inserting, and N as the total number of words.

| Operation | Time Complexity | Explanation |
|------------------|-----------------|---|
| Insertion | $O(L)$ | You strictly process each character of the word once. |
| Search | $O(L)$ | You only traverse the length of the word. Even if the DB has 10 million words, searching for “cat” only takes 3 steps! |
| Space | Variable | In the worst case, it can be large, but it saves space when many words share prefixes (like “tele-” in telephone, telegraph, television). |

Visualization of Efficiency:

```
Searching for "ZOO" in a List of 1 Million Words:  
[Apple, Banana, ... ... ... ... Zoo]  
check, check, ... ... ... ... found  
Steps: 1,000,000 ( $O(N)$ )
```

```
Searching for "ZOO" in a Trie:  
(Root) -> [Z] -> [O] -> [O]  
Steps: 3 ( $O(L)$ )
```

5. Code Implementation

Here is how you build a Trie in Python and JavaScript.

Python Implementation

```
class TrieNode:  
    def __init__(self):  
        # Dictionary to store children nodes  
        # Key: Character, Value: TrieNode  
        self.children = {}  
        # Flag to mark if a word ends at this node  
        self.is_end_of_word = False  
  
class Trie:  
    def __init__(self):  
        self.root = TrieNode()  
  
    def insert(self, word):  
        node = self.root  
        for char in word:  
            # If char not in children, create new node  
            if char not in node.children:  
                node.children[char] = TrieNode()  
            # Move to the child node  
            node = node.children[char]  
        # After loop, mark the current node as end of word  
        node.is_end_of_word = True  
  
    def search(self, word):  
        node = self.root  
        for char in word:  
            if char not in node.children:  
                return False # Path doesn't exist
```

```

        node = node.children[char]
# Return True only if we are at the marked end of a word
return node.is_end_of_word

def starts_with(self, prefix):
    node = self.root
    for char in prefix:
        if char not in node.children:
            return False
        node = node.children[char]
    return True

# Example Usage
trie = Trie()
trie.insert("apple")
print(trie.search("apple"))    # True
print(trie.search("app"))     # False (word not finished)
print(trie.starts_with("app")) # True (prefix exists)

```

JavaScript Implementation

```

class TrieNode {
    constructor() {
        this.children = {};
        this.isEndOfWord = false;
    }
}

class Trie {
    constructor() {
        this.root = new TrieNode();
    }

    insert(word) {
        let node = this.root;
        for (let char of word) {
            if (!node.children[char]) {
                node.children[char] = new TrieNode();
            }
            node = node.children[char];
        }
        node.isEndOfWord = true;
    }

    search(word) {
        let node = this.root;

```

```

        for (let char of word) {
            if (!node.children[char]) {
                return false;
            }
            node = node.children[char];
        }
        return node.isEndOfWord;
    }

    startsWith(prefix) {
        let node = this.root;
        for (let char of prefix) {
            if (!node.children[char]) {
                return false;
            }
            node = node.children[char];
        }
        return true;
    }
}

```

6. Applied Learning: LeetCode Strategy

A classic problem to apply this to is **Word Search II** (Find all words from a dictionary in a 2D board of characters).

The “Aha!” Moment: If you try to search for every single word in the dictionary separately on the board, it takes forever. Instead, put all dictionary words into a **Trie**.

How to think about it:

1. **Preprocessing:** Build the Trie from the list of words.
2. **Traversal:** Walk through the 2D grid (using DFS/Backtracking).
3. **Synchronization:** As you move on the grid (e.g., you step on ‘A’, then ‘P’), move the pointer in the Trie simultaneously.
4. **Pruning:** If the Trie pointer says “No word starts with ‘QZ’”, stop searching that path on the grid immediately!

Conceptual Diagram for Grid Search:

| | |
|------------|--|
| Grid: | Trie Path Check: |
| [C][A] | Start at 'C' -> Exists in Trie? Yes. |
| [T][R] | Move to 'A' -> Exists in Trie? Yes. Move to 'T' -> Exists in Trie? Yes. End of word? Yes! Found "CAT". Backtrack to 'A'. Move to 'R' -> Exists in Trie? Yes. End of word? Yes! Found "CAR". |

7. Comparison with Other Algorithms

| Feature | Trie | Hash Map (HashTable) |
|----------------------|---|---|
| Search Time | $O(L)$ - Fast and consistent | $O(L)$ - Theoretically fast, but collisions can slow it down. |
| Prefix Search | Excellent (Designed for this) | Poor (Must scan all keys or convert to list) |
| Memory Usage | Efficient if many words share prefixes. Expensive if words are very unique (lots of pointers). | Efficient, but stores full strings repeatedly. |
| Ordering | Can print words in alphabetical order easily. | Unordered. |